
RDT Protocol

Release 1.0

Geoffrey Tse

Oct 31, 2025

CONTENTS:

1	HW3	3
1.1	client module	3
1.2	emulator module	3
1.3	protocol module	5
1.4	rdt module	6
1.5	server module	7
	Python Module Index	9

Add your content using reStructuredText syntax. See the [reStructuredText](#) documentation for details.

1.1 client module

client.py

RDT client for GET/PUT via the emulator.

This client sends application-layer commands over a reliable data transfer session (RDTSession), which itself runs over UDP. The packets are routed through the network emulator, allowing tests of loss, corruption, duplication, and reordering.

Usage examples:

```
python client.py --server 127.0.0.1:12000 --emulator 127.0.0.1:10000 GET test.txt python client.py --server 127.0.0.1:12000 --emulator 127.0.0.1:10000 PUT sample.bin
```

Behavior summary:

- GET <file> - request a file from server, store locally as download_<file>
- PUT <file> - upload a local file to the server

client.main()

Entry point for the RDT file client.

Steps:

1. Parse command-line arguments.
2. Bind a local UDP socket.
3. Create an RDTSession that sends to the emulator (NOT directly to server).
4. Run a background thread to pump raw UDP packets into the session.
5. Execute either GET or PUT.
6. Shutdown cleanly, ensuring retransmission thread exits before closing socket.

client.win_udp_no_connreset(sock: socket) → None

Disable WSAECONNRESET behavior on Windows UDP sockets.

When a UDP datagram is sent to a port where nothing is listening, Windows may raise WinError 10054 on recv calls. This function disables that behavior so the client does not crash if the emulator/server closes.

1.2 emulator module

emulator.py

UDP packet emulator to inject loss, corruption, duplication, and reordering. Works as a simple A<->Server relay.

1.2.1 Module overview

This tool sits between a client (A) and a real server, forwarding UDP datagrams in both directions while optionally introducing network impairments. It is used to test the robustness of a reliable data transfer protocol (e.g., selective retransmission, checksums, and in-order delivery), by simulating imperfect network conditions:

- Loss: drop selected packets
- Corruption: flip a random byte inside a packet
- Duplication: send the same packet twice
- Reordering: delay one packet briefly to reorder relative to others

`emulator.corrupt(data: bytes) → bytes`

Return a copy of ‘data’ with one randomly chosen byte bit-flipped (XOR 0xFF).

This simulates payload corruption on the wire.

Parameters

data – Original datagram bytes.

Returns

Corrupted datagram bytes (or unchanged if empty).

`emulator.emulator_loop(sock_a, addr_a, sock_b, addr_b, server_addr, args, stop_evt: Event)`

Main forwarding loop that applies impairments and relays packets between A and the server.

Parameters

- **sock_a** – UDP socket bound to the emulator’s “A” listening address.
- **addr_a** – Tuple (ip, port) for the emulator “A” side (for logging).
- **sock_b** – UDP socket bound to the emulator’s “B” listening address.
- **addr_b** – Tuple (ip, port) for the emulator “B” side (for logging).
- **server_addr** – Tuple (ip, port) of the actual server destination.
- **args** – Parsed argparse Namespace containing impairments probabilities.
- **stop_evt** – Event used to request graceful termination (e.g., on SIGINT).

Notes

- Sockets are set non-blocking and monitored with `select.select()`.
- All send/recv operations are wrapped to prevent unexpected exceptions from exiting the loop; errors are logged and the loop continues.

`emulator.main()`

Parse CLI arguments, create/bind emulator sockets, and run the emulator loop.

Expected usage:

```
python emulator.py --listen-a 127.0.0.1:10000 --listen-b 127.0.0.1:10001
--server 127.0.0.1:12000 --loss 0.1 --corrupt 0.05 --reorder 0.05 --dup 0.02
```

Parameters

- **--listen-a** – IP:port where the emulator listens for the client (“A”) side.
- **--listen-b** – IP:port where the emulator listens for the server-facing side.

- **--server** – IP:port of the actual server that the emulator forwards to.
- **--loss** – Probability [0..1] to drop a packet.
- **--corrupt** – Probability [0..1] to corrupt a packet.
- **--reorder** – Probability [0..1] to hold-and-delay (reorder) a packet.
- **--dup** – Probability [0..1] to duplicate a packet (send twice).

`emulator.win_udp_no_connreset(sock: socket) → None`

Disable Windows-specific UDP “connection reset” behavior for a socket.

On Windows, when a UDP datagram is sent to a port where nothing is listening, the OS may WinError 10054 on subsequent `recv` calls. This helper disables that behavior.

Parameters

sock – A bound UDP socket.

1.3 protocol module

`protocol.py`

This module defines the packet structure used by the Reliable Data Transfer (RDT) protocol. It provides:

- A fixed binary header format shared by both sender and receiver.
- Support for packing and unpacking packets into raw bytes suitable for UDP.
- A CRC32 checksum for data integrity validation.
- A compact payload limit chosen to enforce low data rate (< 500 bps) when combined with the enforced send delay in *rdt.py*.

class `protocol.Packet(seq_num: int, ack: bool, payload: bytes = b'')`

Bases: `object`

Represents a single RDT protocol packet, containing:

- Sequence number
- ACK flag (data or acknowledgement)
- Variable-length payload (0 to MAX_PAYLOAD bytes)
- CRC32 checksum for full data integrity

The class supports packing/unpacking into raw bytes suitable for UDP transmission.

pack() → `bytes`

Serialize the packet into raw bytes suitable for sending via UDP. CRC32 checksum is computed over the header (without checksum field) + payload.

Returns

Raw bytes that represent the packet on the wire.

static unpack(buf: bytes) → Tuple[*Packet*, bool]

Parse raw bytes into a *Packet* object and verify checksum integrity.

Parameters

buf – Raw UDP bytes received

Returns

(*Packet* object, checksum_valid flag)

Raises**ValueError** – if buffer is too small to contain a header

1.4 rdt module

rdt.py

Selective Repeat reliable data transfer over UDP.

Key points: - Reliability: CRC32, per-packet timers, selective retransmissions - Reordering support with in-order delivery to the app - Very low send rate without congestion control:

We enforce < 500 bps by using small payloads (32B) and a fixed 0.6s gap between packets.

1.4.1 Module overview

This module implements the sender/receiver state machines for a minimal Selective Repeat (SR) protocol over UDP. It relies on *protocol.Packet* for serialization and integrity (CRC32). The *RDTSession* class encapsulates all per-peer state:

- **Sender side:**

- Sliding window (size = DEFAULT_WINDOW)
- Per-packet timers and selective retransmissions
- Throttled send rate (SEND_GAP) to satisfy the sub-500 bps requirement

- **Receiver side:**

- Validates CRC32
- Sends per-packet ACKs
- Buffers out-of-order data; delivers in order to the application via a queue

```
class rdt.RDTSession(sock: socket, peer: Tuple[str, int], window: int = 8, timeout: float = 2.0)
```

Bases: object

Per-peer Selective Repeat session.

handle_raw(raw: bytes)

Demultiplex a raw UDP datagram into ACK or DATA, verify checksum, and dispatch to the appropriate handler.

Parameters**raw** – Raw datagram bytes as received from the socket.**recv_available**() → bytes

Retrieve any in-order bytes available for the application without blocking.

Returns

Concatenation of all currently queued in-order chunks (maybe empty).

Return type

bytes

send(data: bytes)

Reliable send with windowing and fixed inter-packet gap.

Splits ‘data’ into chunks of size at most MAX_PAYLOAD, then:

- 1) Waits until there is space in the sliding window.

- 2) Assigns a new sequence number and packs the chunk into a Packet.
- 3) Sends the packet to the peer and starts its retransmission timer.
- 4) Sleeps for SEND_GAP to keep the overall rate < 500 bps.

Parameters

data – Arbitrary bytes to send reliably to the peer.

stop()

Stop the retransmission thread and release resources owned by the session.

`rdt.now()` → float

Monotonic time helper used for timers (immune to wall-clock changes).

Returns

Current monotonic time in seconds.

Return type

float

1.5 server module

server.py

Single-thread UDP server with per-peer RDT sessions and a tiny file service.

The server exposes a minimal application-layer file transfer protocol on top of the custom reliable transport (RDT-Session). Each remote UDP address gets its own RDT session, allowing multiple independent clients to interact at once.

Supported commands (ASCII, newline terminated):**GET <filename>**

-> server sends file contents back via RDT

PUT <filename>

-> client uploads raw file bytes immediately after

(any other text) -> server responds with “OK: ECHO: <text>”

class server.PeerState

Bases: object

Per-peer application state.

- inbuf: accumulates command-line bytes until ‘

‘ is seen

- mode: None or ‘receiving’ during PUT
- filename: target filename during PUT
- filebuf: bytearray of file content during PUT
- last_data_ts: last time we received any file bytes (for idle cutoff)

on_file_bytes(data: bytes) → None

put_done() → Tuple[str, int, bytes]

Finalize PUT state; return (filename, num_bytes, content).

start_put(filename: str) → None

server.main()

Main server loop.

Responsibilities:

- Bind a UDP socket and listen for incoming datagrams.
- Maintain a dictionary of per-peer RDTSession objects + app state.
- Feed raw packets into the proper session via `handle_raw()`.
- Parse newline-terminated commands from a per-peer buffer.
- While in PUT mode, treat all bytes as file data until idle timeout.
- Shut down cleanly on Ctrl+C.

server.win_udp_no_connreset(sock: socket) → None

Windows-only: disable WSAECONNRESET behavior.

On Windows, if a UDP packet is sent to an unreachable port, the OS may raise `WinError 10054` on subsequent `recv` calls. This helper disables that behavior so the server does not crash when a client exits abruptly.

PYTHON MODULE INDEX

C

client, 3

e

emulator, 3

p

protocol, 5

r

rdt, 6

s

server, 7